

Backpropagation in MLP

Tirtharaj Dash

Dept. of CS & IS and APPCAIR
BITS Pilani, Goa Campus

September 04, 2021

Where we are:

Summary of the previous lecture:

- Deep Feedforward Nets and MLPs
- Forward computation in MLPs
- Role of activation functions
- Re-doing XOR with ReLU activation
- Role of hidden representations

Today, we will know:

- How the error is backpropagated via usage of chain rule in calculus.
- That is, we will derive the steps in the backpropagation procedure.
- This will help us understand the tensor-level operation that I will show in the next lecture.

- We are concerned with MLPs (deep fully-connected feedforward neural networks).
- The parameters of an MLP are its synaptic strengths or interconnection weights.
- There are two main computations in an MLP:
 - Forward: Activations flow from inputs to output layer (leading upto: the calculation loss)
 - Backward: Gradient of the loss (or error) is backpropagated from the output layer towards input via the hidden layers (leading upto: the update of weights)


Introduction II

- The parameters (weights) play a significant role in forward and backward computation.
- The method adopted to update the parameters in a deep network is known as the Backpropagation procedure [1].

[1] D.E. Rumelhart, G. Hinton, R.J. Williams (1986): Learning representations by back-propagating errors. *Nature*, 323(6088), 533–536.

Introduction III

- A portion of a 1-hidden layer MLP showing 1 hidden unit:



Notations I

We will adopt the following for our derivations:

- For convenience, we will consider a 1-hidden layer MLP.
- L denotes the loss function: At this point, we will not define what L is, but, it is some function $L(\mathbf{y}, \hat{\mathbf{y}})$.
- \mathbf{y} : vector of target outputs (y_k denotes the output of k th output neuron).
- $\hat{\mathbf{y}}$: vector of computed outputs
- There are c output neurons (one neuron for each target). For instance:
 - In binary classification, $c = 1$
 - In multiclass classification, $c = \text{no. of classes}$
 - In regression, $c = 1$

- Some examples of $L(\mathbf{y}, \hat{\mathbf{y}})$:
 - Binary cross-entropy: $L = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y}))$
 - Cross-entropy: $L = -\sum_{k=1}^c y_k \log \hat{y}_k$
 - SSE: $L = \frac{1}{2} \sum_{k=1}^c (y_k - \hat{y}_k)^2$ (for a single instance)
 - MSE for a batch of instances: $L = \frac{1}{2m} \sum_{i=1}^m \sum_{k=1}^c (y_k^{(i)} - \hat{y}_k^{(i)})^2$

Notations III

- The subscript i corresponds to the inputs.
- There are d input units, each corresponding to a feature vector in \mathbf{x} .
- The subscript j corresponds to the hidden layer.
- There are m units in the hidden layer.
- The subscript k corresponds to the output layer.
- There are c units in the output layer.

- Weights (see the footnote below):
 - w_{ji} : weight parameter for the connection from i th input to j th neuron in the hidden layer.
 - w_{kj} : weight parameter for the connection from j th hidden neuron to k th neuron in the output layer.

Note: In our forward propagation lecture (also in our Textbook 1), we were using a different notation. That is, w_{ij} and w_{jk} respectively. Don't get confused.

- z denotes the net input to any neuron j . For example,
 - z_j : for the hidden neuron j
 - z_k : for the output neuron k
- a denotes the activation value. That is: $a = \sigma(z)$, where σ is some activation function (preferably, continuous).
 - a_j : for the hidden neuron j
 - a_k : for the output neuron k . But, this is equal to \hat{y}_k .

Gradient Descent (GD) on loss:

- The update to any parameter w_{ij} at any iteration t is carried out using the step

$$w_{ij}^{(t)} \leftarrow w_{ij}^{(t-1)} - \eta \left. \frac{\partial L}{\partial w_{ij}} \right|_{\text{at } w_{ij}=w_{ij}^{(t-1)}}$$

- Next, we will see how to derive $\frac{\partial L}{\partial w_{ij}}$ for the hidden layer parameters and the output layer parameters.

Note: In this particular slide and in the next one the subscript in w (i.e. i and j) are just local variables. They represent any weight in a network.

What exactly is $\frac{\partial L}{\partial w_{ij}}$?

- This is calculated using chain-rule of (partial) derivative.
This turns out to be:

$$\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial \text{activation}} \frac{\partial \text{activation}}{\partial \text{net-input}} \frac{\partial \text{net-input}}{\partial w_{ij}}$$


- or,

$$\frac{\partial L}{\partial w_{ij}} = \frac{\partial L}{\partial a} \frac{\partial a}{\partial z} \frac{\partial z}{\partial w_{ij}}$$

Ignoring the subscript on a and z here. In the next slides, these will be clearer.

GD and Backprop III

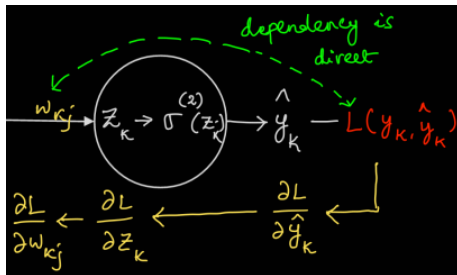
A figure showing flow of gradient from L towards the inputs:



Gradients at output layer I

We want to compute $\frac{\partial L}{\partial w_{kj}}$ for all j and k :

- Let us see the dependency graph (chain) first:



- Now, the expression for the derivative based on the above chain:

$$\frac{\partial L}{\partial w_{kj}} = \frac{\partial L}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial z_k} \frac{\partial z_k}{\partial w_{kj}}$$

- We want to compute these individual quantities. See next.

- Derivative of the loss w.r.t. the output (activation), $\frac{\partial L}{\partial \hat{y}_k}$:
 - Since we have not defined our function L : Let us keep it as it is, and call it: L'
 - If L is SSE, this would be:

$$\frac{\partial L}{\partial w_{kj}} = -(y_k - \hat{y}_k)$$

Gradients at output layer III

- Derivative of the output (activation) w.r.t. the net input, $\frac{\partial \hat{y}_k}{\partial z_k}$:
 - We know that $\hat{y}_k = \sigma^{(2)}(z_k)$, where $\sigma^{(2)}$ is the output layer activation function, we have then:

$$\frac{\partial \hat{y}_k}{\partial z_k} = \frac{\partial \sigma^{(2)}(z_k)}{\partial z_k} = \sigma^{(2)'}(z_k)$$

Gradients at output layer IV

- Derivative of the net input w.r.t. the weight, $\frac{\partial z_k}{\partial w_{kj}}$:
 - We know that $z_k = \sum_{j=1}^m w_{kj} a_j$ (for simplicity we are not using the bias term; however, including this is straightforward: the summation will start at $j = 0$). Now, the derivative is:

$$\frac{\partial \hat{y}_k}{\partial z_k} = a_j$$

Gradients at output layer V


- Substituting the quantities obtained earlier we get our $\frac{\partial L}{\partial w_{kj}}$ expression:

$$\frac{\partial L}{\partial w_{kj}} = L'(\mathbf{y}, \hat{\mathbf{y}}) \sigma^{(2)'}(z_k) a_j$$

- The blue expression is called the *local gradient* computed at the output neuron k , and we denote it as δ_k .

Gradients at output layer VI

- A simple diagram shown as:




$$\delta_k = L'(\mathbf{y}, \hat{\mathbf{y}}) \sigma^{(2)'}(z_k)$$

- Notice that we do not need to re-calculate this δ_k while calculating some other weight (i.e. different j) for the same k . We can re-use the computed δ_k .
- Also, recall that a_j was already computed during our forward computation.

Gradients at hidden layers I

We want to compute $\frac{\partial L}{\partial w_{ji}}$ for all i and j :


- Let us first see the dependency graph:



The dependency graph does not show that the weight w_{ji} leads to error at all the neurons at the next layer (outputs). We show this next.

Gradients at hidden layers II

- All the paths from w_{ji} to L :



The yellow-lines show the paths.

Gradients at hidden layers III

- What is this really?
 - L depends on $\hat{\mathbf{y}}$
 - $\hat{\mathbf{y}}$ depends on a_j
- For the hidden unit j , we know these following from forward computation:

$$z_j = \sum_{i=1}^d w_{ji} x_i$$

and

$$a_j = \sigma^{(1)}(z_j)$$

- For the output unit k , we know that:

$$z_k = \sum_{j=1}^m w_{kj} a_j$$

Gradients at hidden layers IV

- In the end, $\frac{\partial L}{\partial w_{ji}}$ turns out to be:

$$\begin{aligned}\frac{\partial L}{\partial w_{ji}} &= \left(\sum_{k=1}^c \frac{\partial L}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial z_k} \frac{\partial z_k}{\partial a_j} \right) \frac{\partial a_j}{\partial z_j} \frac{\partial z_j}{\partial w_{ji}} \\&= \left(\sum_{k=1}^c \frac{\partial L}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial z_k} w_{kj} \right) \sigma^{(1)'}(z_j) x_i \\&= \left(\sum_{k=1}^c L'(\mathbf{y}, \hat{\mathbf{y}}) \sigma^{(2)'}(z_k) w_{kj} \right) \sigma^{(1)'}(z_j) x_i \\&= \left(\sum_{k=1}^c \delta_k w_{kj} \right) \sigma^{(1)'}(z_j) x_i \\&= \delta_j x_i\end{aligned}$$

Gradients at hidden layers V

- δ_j is the local gradient computed at the hidden unit j :

$$\delta_j = \left(\sum_{k=1}^c \delta_k w_{kj} \right) \sigma^{(1)'}(z_j)$$

- Notice the first term in parenthesis. This is saying weigh all the local gradients computed at the output layer by the corresponding interconnection weights.

Generalisation to Deeper Nets I

More on δ_j computed in the last slide:

- This is a beautiful finding, which takes all our worries of computing gradients for deeper network.
- This allows us to generalise to any layers, since we are travelling from L towards the inputs, on the way, we are calculating these local gradients.

Generalisation to Deeper Nets II

- Then the local gradient of any neuron j at any hidden layer ℓ (where $\ell_{max} > \ell \geq 1$) can be easily computed as:

$$\delta_j^{(\ell)} = \left(\sum_{k=1}^{m^{(\ell+1)}} \delta_k^{(\ell+1)} w_{kj}^{(\ell+1)} \right) \sigma^{(\ell)'}(z_j^{(\ell)})$$


- Then the gradient of the loss L w.r.t to any weight at layer ℓ is simply:

$$\frac{\partial L}{\partial w_{ji}^{(\ell)}} = \delta_j^{(\ell)} a_i^{(\ell-1)}$$

Notations: (ℓ) is used to denote a layer index; $\ell = 0$ corresponds to the inputs; ℓ_{max} corresponds to the output layer; $m^{(\ell)}$ denotes the number of neurons (units) at layer ℓ ; $\mathbf{a}^{(0)} = \mathbf{x}$.

Generalisation to Deeper Nets III

- This can be clearly shown as a bunch of computational (flow) paths in the network:



Parameter updates I

- We have now computed the derivatives of L w.r.t. all the parameters in the network.
- The update to the parameters based on the GD procedure for any network layer ℓ is then:

$$w_{ij}^{(\ell)} = w_{ij}^{(\ell)} - \eta \frac{\partial L}{\partial w_{ij}^{(\ell)}}$$

- That is, the gradient of L w.r.t. all the parameters of the network have to be computed **first** and then the update is done. **Otherwise** the computation of the local gradients will be **wrong**.

Parameter updates II

- Common mistake:

For layer $\ell \in \{1, \dots, \ell_{\max}\}$:

 Compute gradient of L w.r.t. weights at layer ℓ

 Update the weights at layer ℓ

- Correct (writing again):

For layer $\ell \in \{1, \dots, \ell_{\max}\}$:

 Compute gradient of L w.r.t. weights at layer ℓ

For layer $\ell \in \{1, \dots, \ell_{\max}\}$:

 Update the weights at layer ℓ

Parameter updates III

- The derivations so far has been to show how the backpropagation procedure operates at the scalar-level.
- However, you will notice that many computations such as the local gradients at neurons at any layer can be computed in parallel (i.e. computation of $\delta_i^{(\ell)}$ and $\delta_j^{(\ell)}$ is independent) allowing tensor-level computation.
- Further, the quantities computed at the forward propagation such as a_j s can be reused for backpropagation. This requires some kind of *recor keeping*. This results in requirement of more memory storage for backpropagation.

These two questions will make you comfortable with the derivations we did here:

- Q1 Derive the complete backpropagation procedure for a classification problem with c classes. Here, L is the cross-entropy loss.
- Q2 Derive the complete backpropagation procedure for a regression problem. Here L is the SSE.
- Q3 What happens to the derivations if we have a batch-size greater than 1 in our training? What would change in our derivations?