Representation Learning

Tirtharaj Dash

Dept. of CS & IS and APPCAIR
BITS Pilani, Goa Campus

November 23, 2021

Representation learning :-

Learning (compact) numeric representations for inputs.

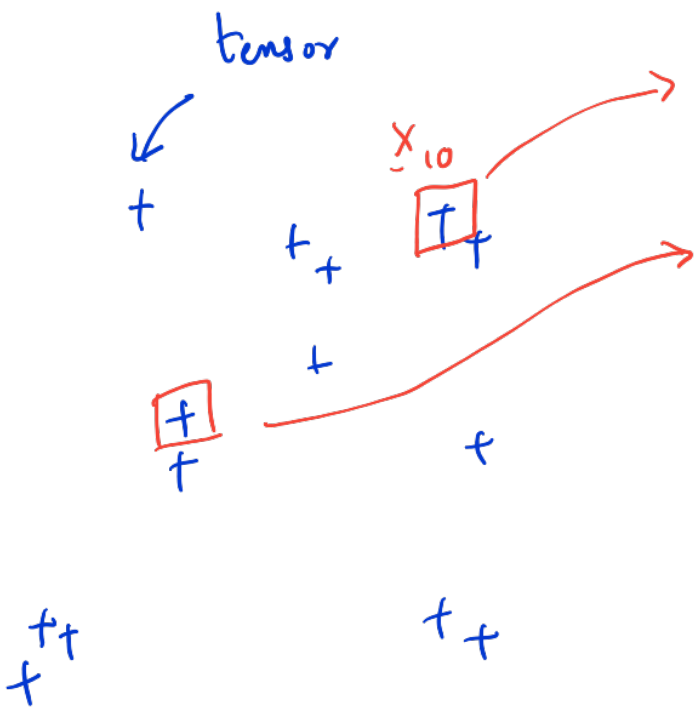

$\underline{z} = W \cdot \underline{x}$

$\underline{h} = g(\underline{z})$

$\underline{h}_{10}$ ← representation for the data instance $\underline{X}_{10}$

HL

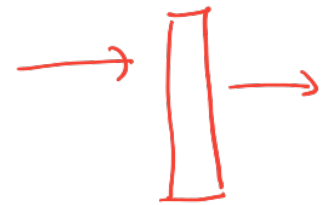

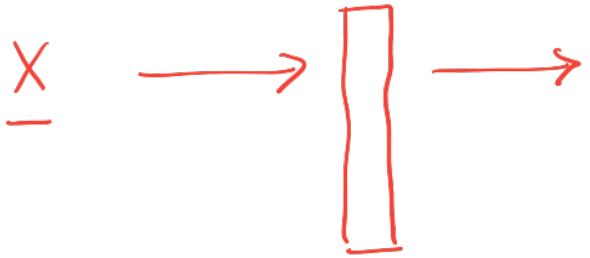

Unsupervised learning :-

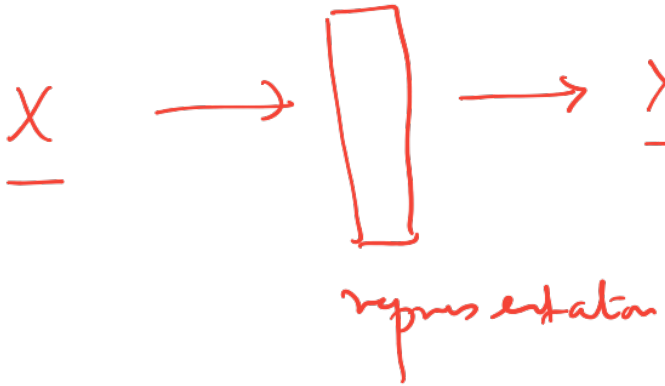

(1) Autoencoder: It learns $p(\underline{h} \mid \underline{x})$.

given some $\underline{x}$ it produces/constructs a code

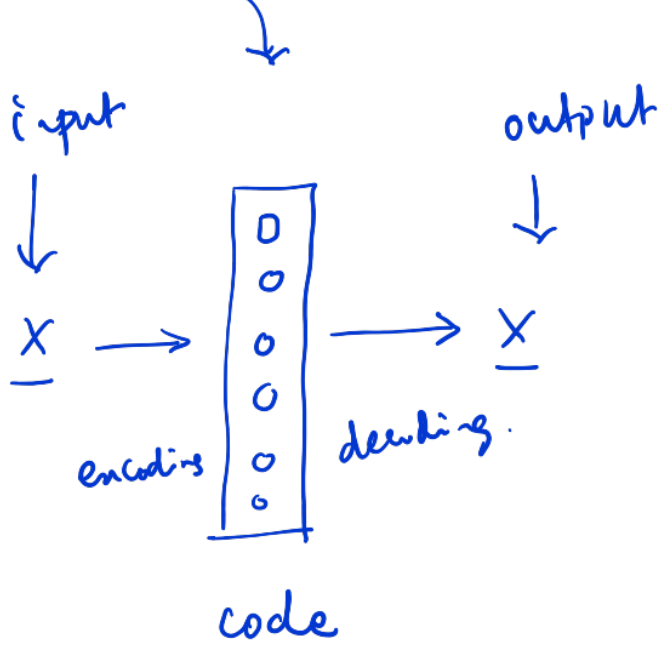

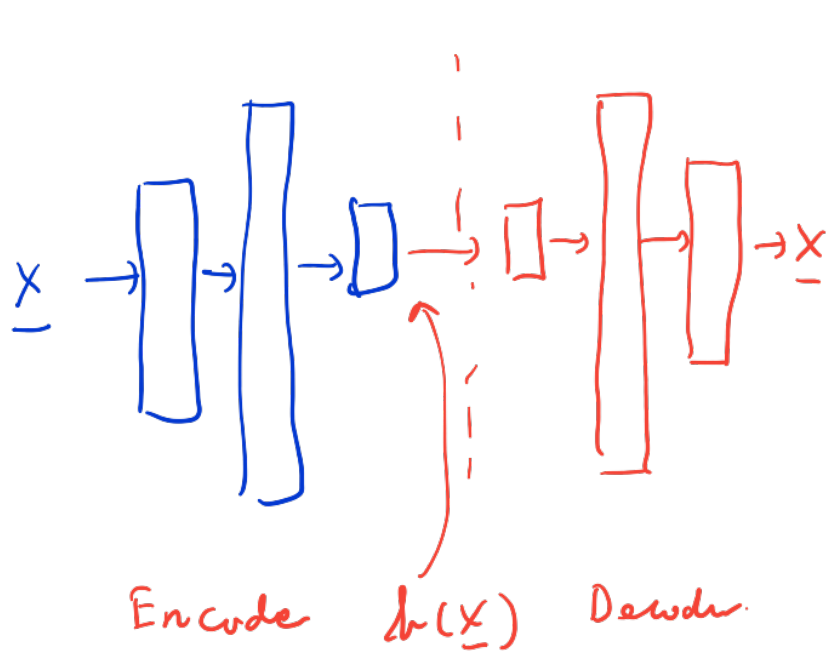

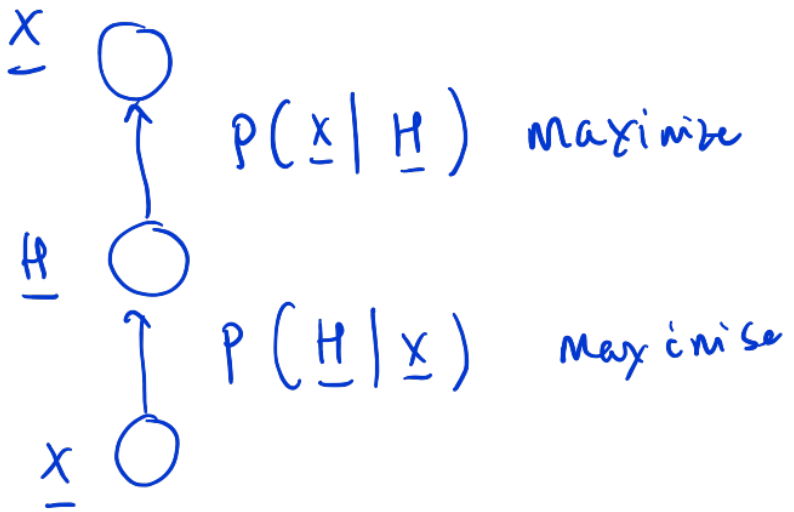

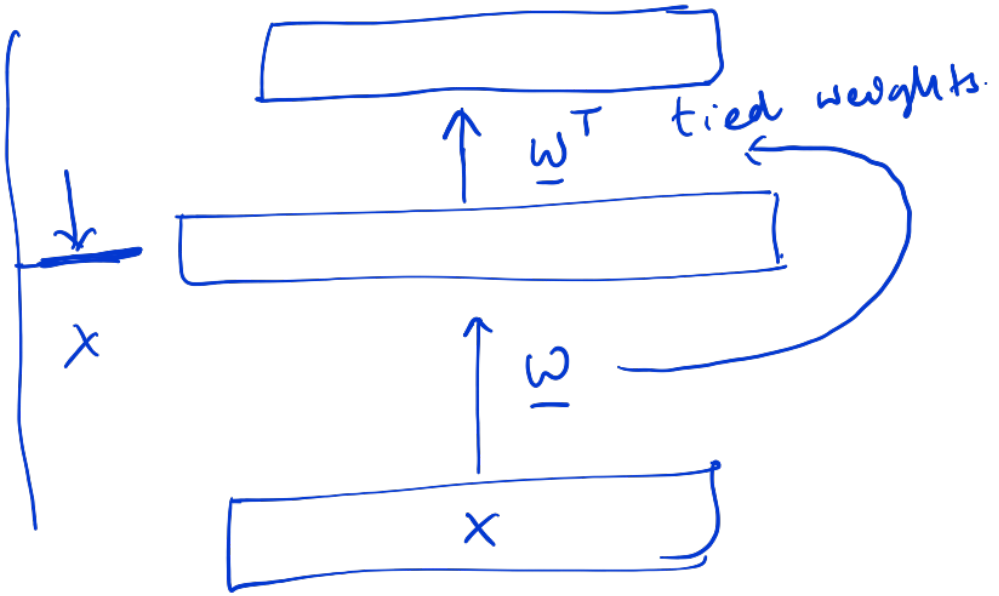

$$\underline{x} \longrightarrow h(\underline{x}) \longrightarrow \hat{\underline{x}}$$

reconstructed data instance

Loss function ?

① MSE : $\underline{x} \in \mathbb{R}^d$

② NLL : ? $\underline{x} \in \{0,1\}^d$

Loss : $L(\underline{x}, \hat{\underline{x}})$ reconstruction error.

Questions:

(1) What if code is sparse ? → length(h(x)) > length($\underline{x}$)

whether the code should be dense !

MNIST digits :- handwritten Digits 0-9.

28 x 28

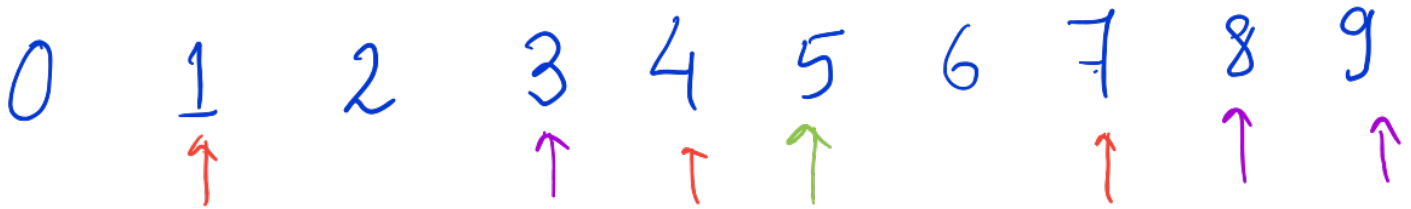

AE for MNIST.

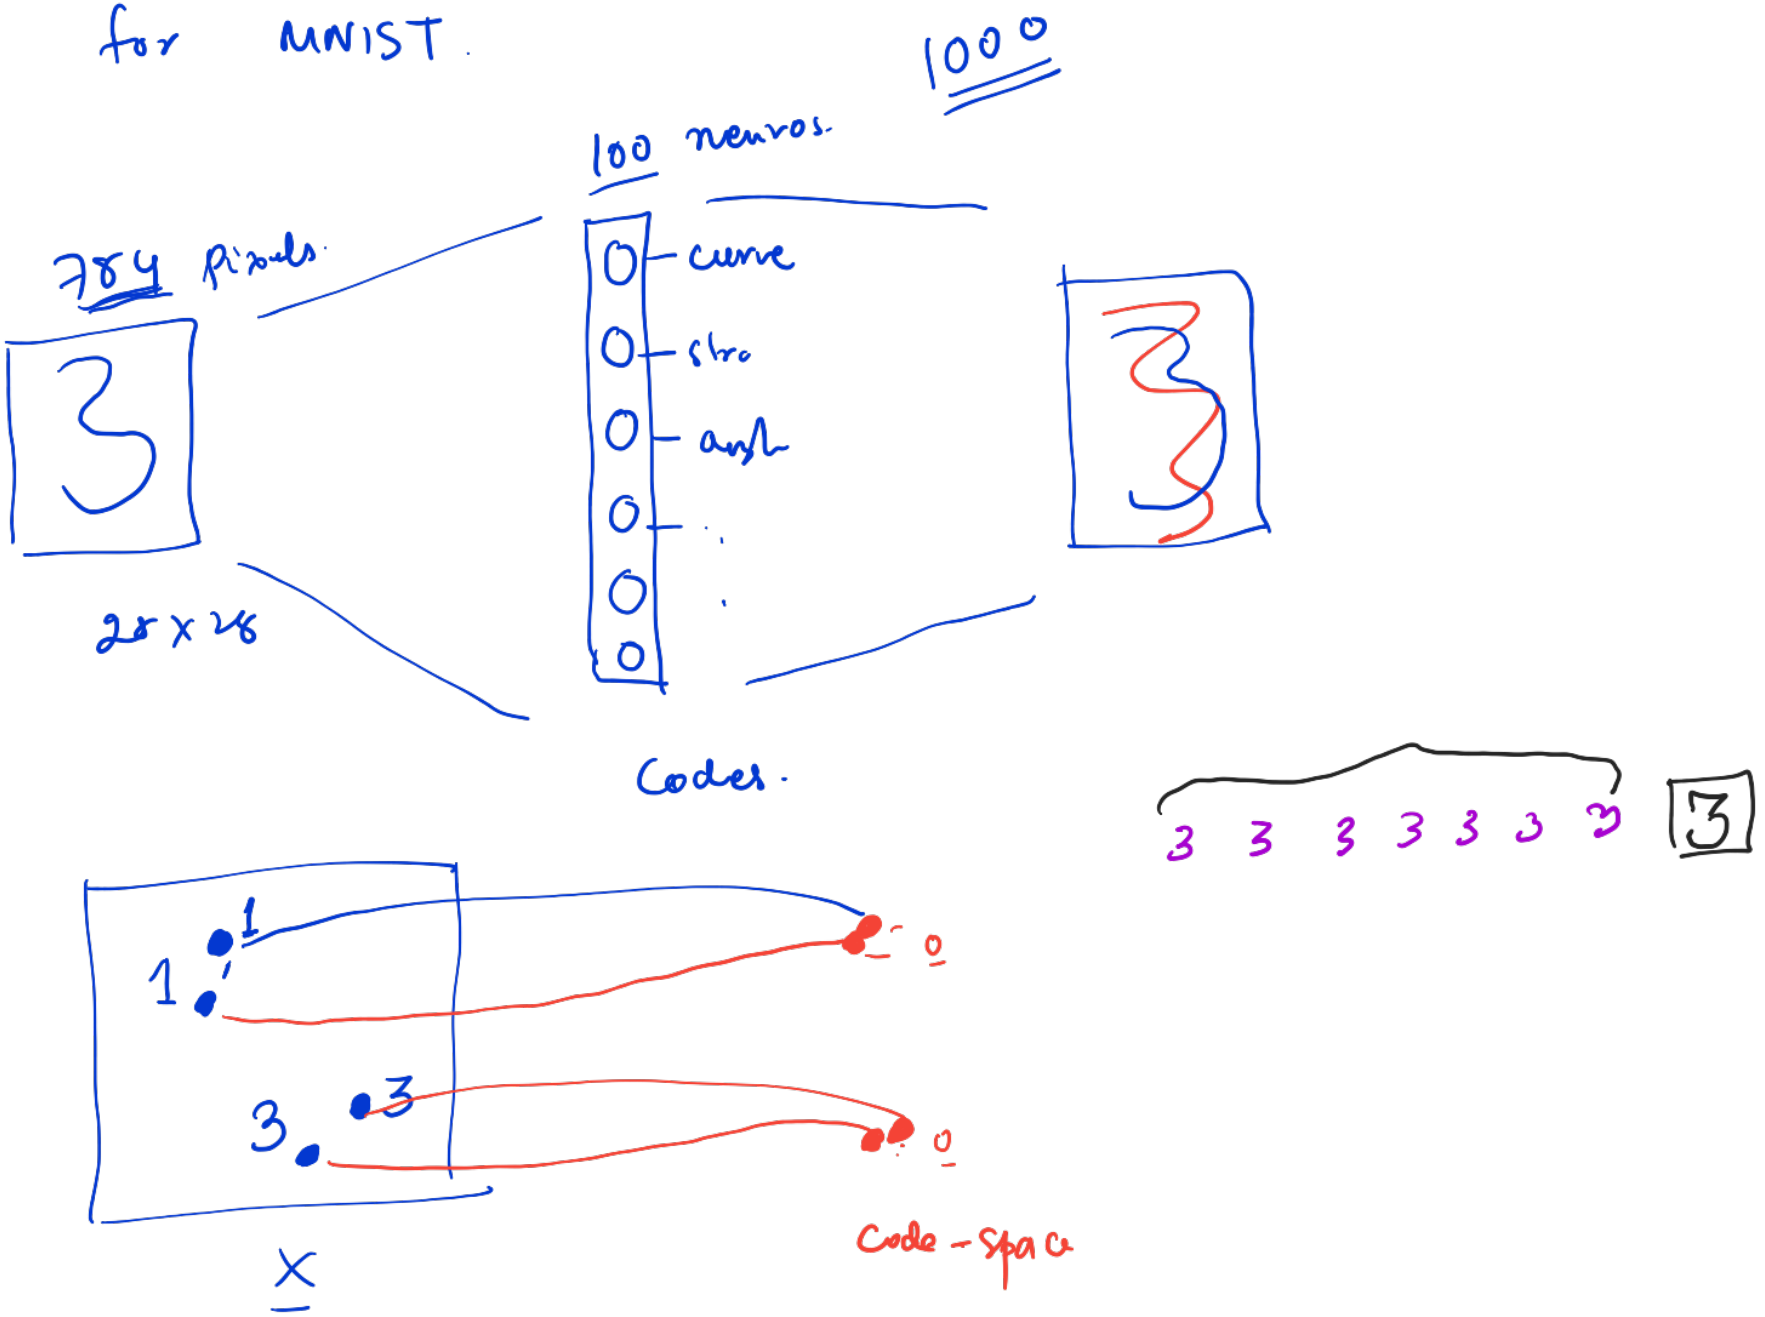
1000
100 neurons.
784 pixels.
28 x 28
curve
slo
ank
Codes.
3 3 3 3 3 3 3
3
1
1
0
3
3
0
X
Code-Space

How to enforce that the code should be "compact"?

Reconstruction error: $L(\underline{x}, \hat{\underline{x}})$

$+ \lambda$ regularisation.

sparsity penalty.

↑

Sparsity regularisation

top-k features.

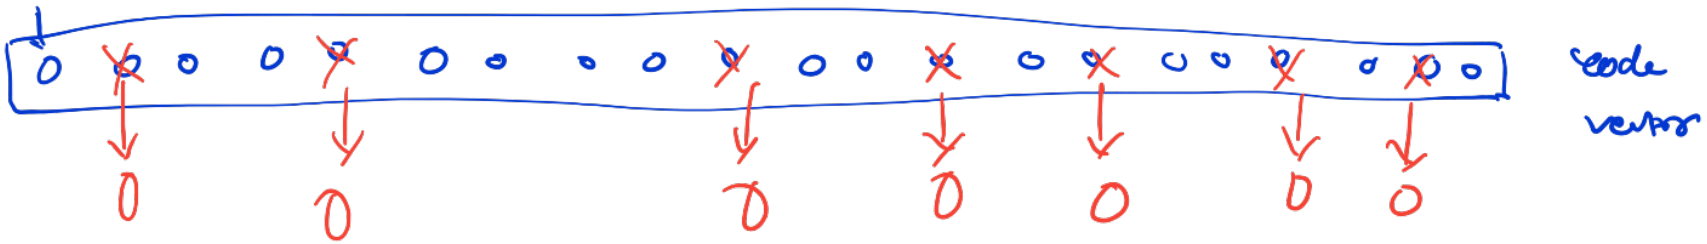

→ penalising the hidden vector with L_1 or L_2.

(Sparse AE).

Problem: What if the data is noisy?

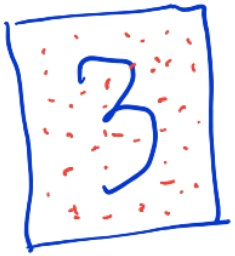

Assume: that the AE was trained with clean data.

Noisy input

X + noise

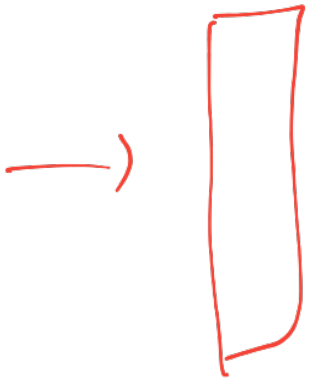

$\hat{X}$

Denoising AE.

1 → provide as input $x + \text{Noise}$

2 → learn the code

3 → try to reconstruct the clean x.

Noise :

① Gaussian noise : $\sigma \sim \mathcal{N}(0, \lambda)$ (λ: regularisation parameter)

(real-valued inputs)

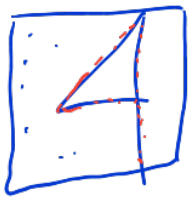

② Masking noise : set a fraction of inputs to 0

(Binary inputs)

③ Salt-pepper noise : set a fraction of inputs to max or min values.

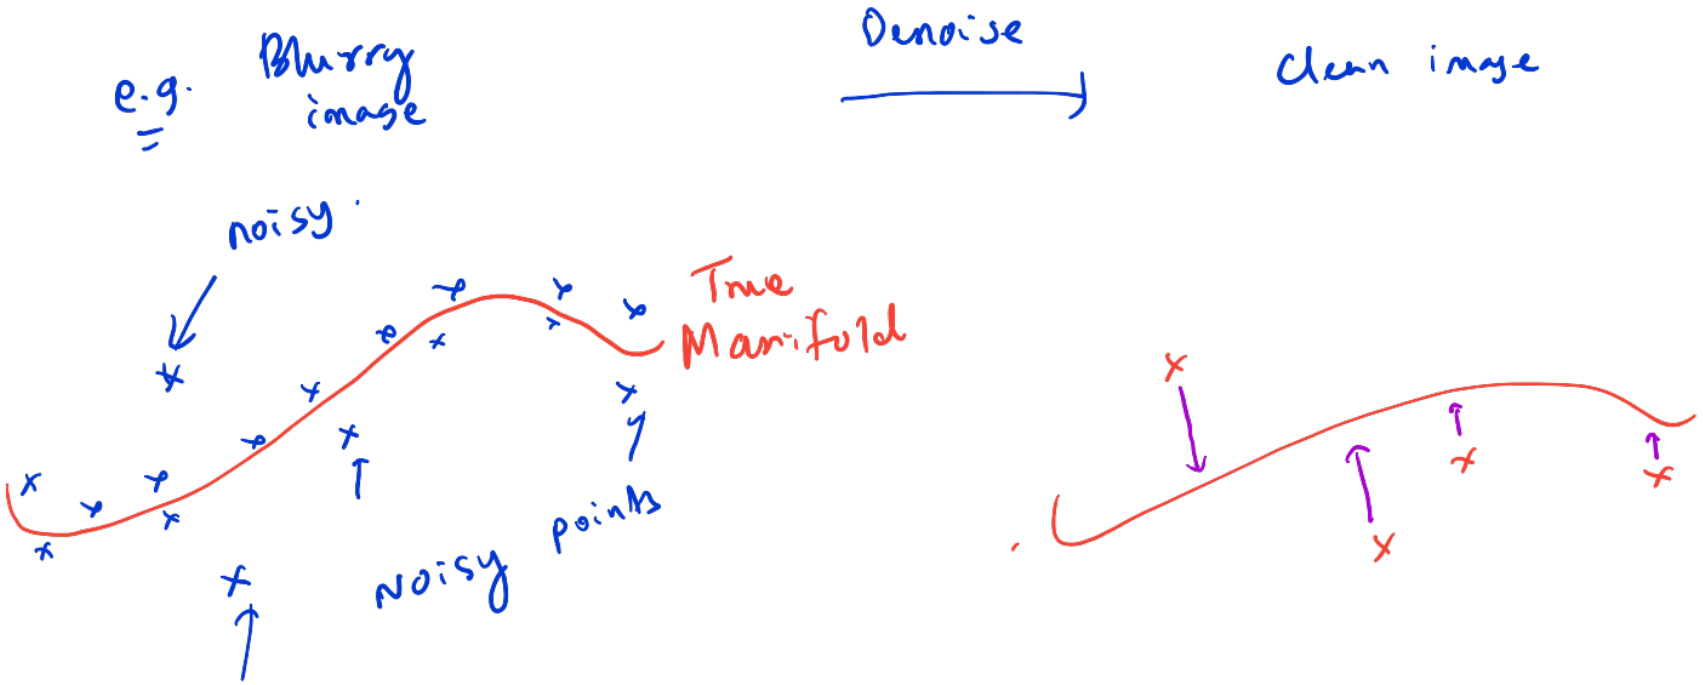
e.g.
Blurry image
Denoise
clean image
noisy
True Manifold
noisy points

Contractive autoencoding :

auto encoders.

AE. : Discriminative model

$P(H|x)$ $P(x|H)$