

9th of Dec

Generative Adversarial Networks (GANs)

Tirtharaj Dash

Dept. of CS & IS and APPCAIR
BITS Pilani, Goa Campus

December 07, 2021

$\mathcal{H}(-, -)$
M, σ

AE
ENC $\begin{matrix} \text{Generator} \\ \text{M} \\ \text{-DEC} \\ \sigma \end{matrix}$

VAE

→ Disentangled
VAE

$p(z)$
↓

→ β -loss
(β -VAE)

→ Vector-Quantised
VAE.
(VQ-VAE)

Main Idea

- GAN consists of two different models:

- **Discriminative model:** estimates the conditional prob. $p(y|x)$ of the label y , given the data instance x . (e.g. MLP, CNN)

- **Generative model:** estimates the joint prob. $p(x, y)$ or $p(x)$

- Note: Discriminative model can only be used in an supervised setting; whereas geneative model can be used in both supervised as well as in unsupervised setting.
- Example: In VAE, a Gaussian distro is used to generate sample and passed to a decoder to (re-)construct the original data instance.

Main Idea




- Generative model inside a GAN:

- It produces synthetic examples of objects that are simple to a real repository of examples.
- The goal is to create synthetic objects that are so realistic that it is hard for a trained observer to say that it is synthetic.

dataset

Main Idea



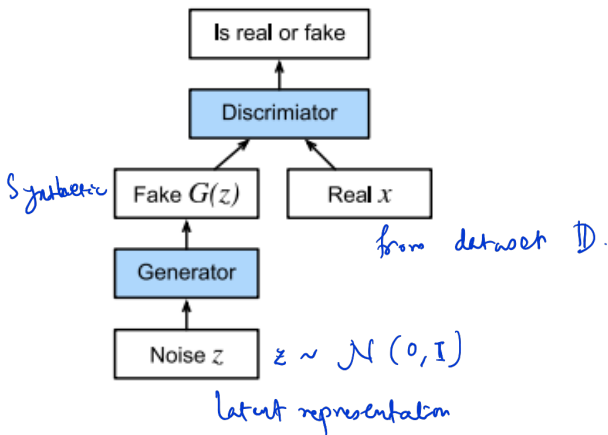
- Discriminative model inside GAN:

- It takes in inputs of either real examples from the base data or synthetic objects created by generator and tries to discriminate between them.
- The goal is to correctly classify the real and synthetic objects.

- Why “adversarial”? The goals of the generator and the discriminator are quite opposite, i.e. the generator wants to fool the discriminator which does not want to be fooled.

Main Idea

GAN architecture:



Training a GAN

Discriminator:

real or synthetic (binary clf.)

- Let \mathbf{x} be any data instance. The discriminator is a binary classifier to distinguish if \mathbf{x} is real (for real data) or synthetic (from generator).
- This can be written as:

$p(y|\mathbf{x})$
 $\in \mathbb{D} \quad G(z)$

Discriminator function $D(\mathbf{x}; \underline{w})$

$$D(\mathbf{x}) = \frac{1}{1 + \exp(-o)} \quad \leftarrow \text{sigmoid}(o)$$

for some output $o \in \mathbb{R}$ computed at the output neuron of the discriminator.

- Let the label y be 1 for real data and 0 for synthetic data.
- The training minimises the cross-entropy i.e.

$$\min_{\underline{D}} \{ -y \log D(\mathbf{x}) - (1 - y) \log(1 - D(\mathbf{x})) \}$$

$p(y|\mathbf{x})$

Images:
 $D: \text{CNN}$

Training a GAN


latent representation
 $\mathbf{z} \sim \mathcal{N}(0, \mathbf{I})$ \longrightarrow synthetic data instance
(noise)

Generator:

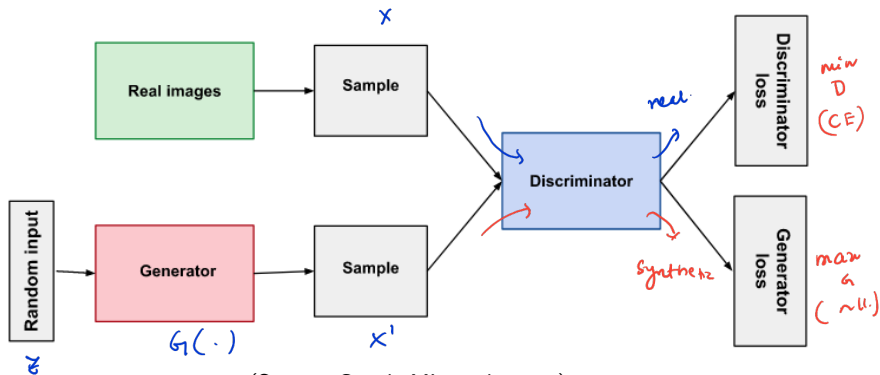
$h \times w \times d$

- It first draws some parameter $\mathbf{z} \in \mathbb{R}^d$ from some some distro. e.g. $\mathbf{z} \sim \mathcal{N}(0, 1)$.
- It then applies a function on \mathbf{z} to generate an instance: $\mathbf{x}' = \underline{G}(\mathbf{z})$ (like a decoder network) *generator*
- The goal of the generator is to produce \mathbf{x}' such that $\boxed{D(\mathbf{x}') \approx 1}$.
- That is: For a given discriminator \underline{D} , we update the parameters of the generator G to maximise the cross-entropy loss when $\underline{y} = 0$, i.e.

$$\max_G \{ \underbrace{-(1-y) \log(1 - D(G(\mathbf{z})))}_{-(1-0) \log(1 - D(\mathbf{x}'))} \} = \max_G \{ \underbrace{-\log(1 - D(G(\mathbf{z})))}_{\mathbf{x}'} \}$$



Training a GAN



(Source: Google ML crashcourse)


Training a GAN

Summary:

- The discriminator and the generator are playing a “minimax” game with one comprehensive objective function


$$\min_{\underline{D}} \max_{\underline{G}} \{ \underbrace{-E_{\mathbf{x} \sim \mathcal{X}} \log D(\mathbf{x})}_{y=1} - \underbrace{E_{\mathbf{z} \sim \mathcal{N}(0,1)} \log(1 - D(G(\mathbf{z})))}_{\text{noise}} \}$$

- The solution is a saddle point (Can you see this?) $y=0$



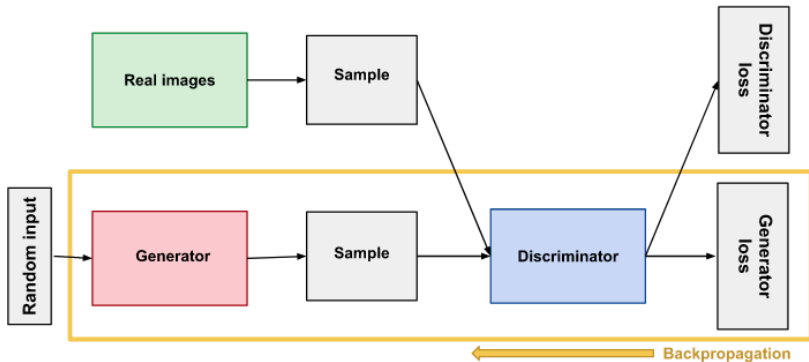
Training a GAN

Backprop:



(Source: Google ML crashcourse)


Training a GAN



(Source: Google ML crashcourse)

Training a GAN

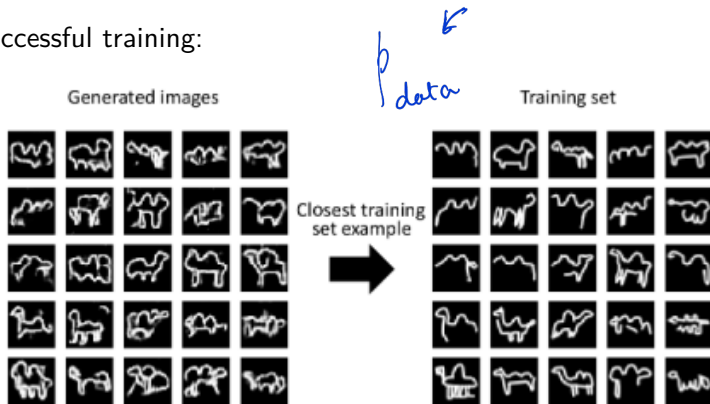
Generated samples at some specific epochs:



(Source: David Foster, Generative Deep Learning)

Training a GAN

With successful training:




(Source: David Foster, Generative Deep Learning)

Known challenges in GAN Training

Oscillating Loss:


- The loss of the discriminator and generator can start to oscillate wildly, rather than exhibiting long-term stability. unstable
- Typically, there is some small oscillation of the loss between batches. However, in the long term the loss should be stable (or gradually decreases or increases). See below:



(Source: David Foster, Generative Deep Learning)

Known challenges in GAN Training

- But, most of the time, the loss of the discriminator and generator start to spiral out. See below:



(Source: David Foster, Generative Deep Learning)

- Also, it is difficult to establish if or when this might occur during training.
- This issue is very common in vanilla GANs.

Known challenges in GAN Training

Mode Collapse:

local minimum

- Mode collapse occurs when the generator finds a small number of samples that fool the discriminator.
- It isn't able to produce any examples other than this limited set.

Known challenges in GAN Training

- How does this occur?
 - Suppose we train the generator over several batches without updating the discriminator in between.
 - The generator would be inclined to find a single observation (also known as a mode) that always fools the discriminator and would start to map every point in the latent input space (z) to this observation.
 - This means that the gradient of the loss function collapses to near 0.

$z \longrightarrow \text{mode}$

Known challenges in GAN Training

- Can retraining the discriminator help?
 - It maynot: The generator will simply find another mode that fools the discriminator.
 - That is, the generator has become numb to the inputs and therefore does not diversify the generation (see below: generator generates identical samples even though its inputs are different).

z_1



z_2




z_3



(Source: David Foster, Generative Deep Learning)

Known challenges in GAN Training




Uninformative Loss:

- It is natural to think that smaller the loss function of the generator, the better the quality of the images produced.
- However, since the generator is only graded against the current discriminator and the discriminator is constantly improving, we cannot compare the loss function evaluated at different points in the training process.

Known challenges in GAN Training

- See below: The loss function of the generator actually increases over time, even though the quality of the images is clearly improving.



(Source: David Foster, Generative Deep Learning)

- This lack of correlation between the generator loss and image quality sometimes makes GAN training difficult to monitor.

Known challenges in GAN Training

Hyperparameter sensitivity:

- GANs involve a huge set of hyperparameters: generator and discriminator structure, training hyperparameters, etc.
- GANs are highly sensitive to very slight changes in all of these parameters, and finding a set of parameters that works is often a case of educated trial and error, rather than following an established set of guidelines.

Known challenges in GAN Training

Some of the above challenges (mainly, mode collapse) can be handled with Wasserstein GAN: <https://arxiv.org/abs/1701.07875>

Wasserstein GAN - Gradient Penalty.

Q1 Generate digits. (dataset: MNIST)

Q2 Generate 10 images of the digit ["]7["]
"3"

Conditional Generation of Samples

↓
Controlled generation of samples.

- Vanilla GANs do not allow controlled generation of samples
 - Example: Generate samples of a specific class
- How can we allow the condition to be incorporated into GAN modelling?


Vanilla GAN

$$D: p(y|x) \quad ; \quad x \rightarrow \hat{y} \quad \text{real/fake}$$
$$\underline{G}: p(x) \quad ; \quad \boxed{\tilde{x}} \rightarrow \underline{x'} \quad \text{y}$$

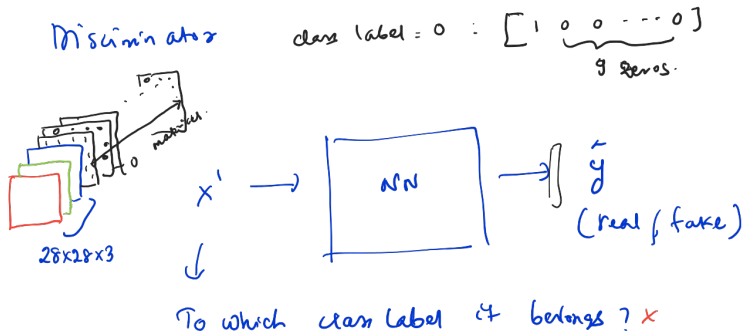
Conditional Generation of Samples

Generator

Conditional generation



Conditional Generation of Samples



10 digits.

